

Parallel Orthogonal Recursive Bisection

Team Metropolis:
James Farzi, JJ Lay, Graham West

April 10, 2019

1 Introduction

2 Implementation

- `main`
- KD tree
 - Building the tree
 - Searching the tree
- Parallel sorting
 - Changes
 - `adaptBins`
- Searching the tree

3 Validation

- MATLAB visualisations
- Other tests

4 Results

5 Conclusions

Introduction

Implementation

Building the tree

To build the tree, we use several functions which perform different aspects/sections of the task

Functions:

- `buildTree`
- `buildTree_serial`
- `buildTree_parallel`
- `getSortDim`

Building the tree

`buildTree` checks the number of compute nodes in the current communicator and determines whether to call the parallel or serial versions of the code

Algorithm 1: `buildTree(...)`

```
1:  $q$  = Size of current communicator
2: if  $q > 1$  then
3:   buildTree_parallel(...)
4: else
5:   buildTree_serial(...)
6: end if
```

Building the tree

`buildTree_serial` performs ORB using a single compute node

Algorithm 2: `buildTree_serial(data, tree, ...)`

```
1: if tree.n > 1 then
2:   Calculate x, y, z mins, maxs, ranges, and partition center
3:   Sort data over sortDim =  $\text{argmax}(x, y, z \text{ ranges})$ 
4:   Split data: dataL, dataR
5:   if  $|dataL| > 0$  then
6:     Create tree.L
7:     buildTree_serial( dataL, tree.L, ... )
8:   end if
9:   if  $|dataR| > 0$  then
10:    Create tree.R
11:    buildTree_serial( dataR, tree.R, ... )
12:   end if
13: else
14:   Store data (a single point)
15: end if
```

Building the tree

`buildTree_parallel` performs ORB using a multiple compute nodes

Algorithm 3: `buildTree_parallel(data, tree, comm, ...)`

```
1: Call getSortDim(...): calculates  $x, y, z$  mins, maxs, ranges, partition center, and  
   returns sortDim  
2: Sort data over sortDim using parallelSort(data, sortDim, comm, ...)  
3: if myRank < numNodes/2 then  
4:   Create tree.L, commL  
5:   buildTree_parallel(data, tree.L, comm, ...)  
6: else  
7:   Create tree.R, commR  
8:   buildTree_parallel(data, tree.R, comm, ...)  
9: end if
```

It is assumed that $tree.n > 1$ will never occur in `build/tree_parallel` since we usually deal with large amounts of data

Building the tree

`getSortDim` finds the longest axis and stores several key tree fields

Algorithm 4: `getSortDim(data, tree, comm, ...)`

- 1: Each process gets its local x, y, z min and max
 - 2: Rank 0 receives these, determines the global x, y, z min and max, determines the `sortDim`, and Bcast's all of these values back to the other nodes
 - 3: The global mins/maxs, partition center, and partition radius are stored in *tree*
 - 4: return *sortDim*
-

serial 501

review

conversion to function rank 0 multiple communicators

old new alternate

Validation

MATLAB Demos:

- 2D
- 3D

Results

Conclusions