

Parallel Orthogonal Recursive Bisection

Team Metropolis:
James Farzi, JJ Lay, Graham West

April 10, 2019

Outline

1 Introduction

2 Implementation

- main
- KD tree
 - Building the tree
 - Searching the tree
- Parallel sorting
 - Changes
 - adaptBins

3 Validation

- MATLAB visualitations
- Other tests

4 Results

5 Conclusions

Introduction

Implementation

Building the tree

Our `main` was quite simple due to our organization of the project into many levels of functions

We also were able to use much of the basic initialization and data importing functions from the previous project

Algorithm 1: `main(...)`

- 1: Initialize MPI
 - 2: Set number of files, lines per file to read
 - 3: import the *data*
 - 4: Initialize *tree*
 - 5: `buildTree(data, tree, comm, ...)`
 - 6: Search the tree with `search501(tree, ...)`
 - 7: Finalize MPI
-

Building the tree

To build the tree, we use several functions which perform different aspects/sections of the task

Functions:

- `buildTree`
- `buildTree_serial`
- `buildTree_parallel`
- `getSortDim`

Building the tree

`buildTree` checks the number of compute nodes in the current communicator and determines whether to call the parallel or serial versions of the code

Algorithm 2: `buildTree(...)`

```
1:  $q$  = Size of current communicator
2: if  $q > 1$  then
3:   buildTree_parallel(...)
4: else
5:   buildTree_serial(...)
6: end if
```

Building the tree

`buildTree_serial` performs ORB using a single compute node

Algorithm 3: `buildTree_serial(data, tree, ...)`

```
1: if tree.n > 1 then
2:   Calculate x, y, z mins, maxs, ranges, and partition center
3:   Sort data over sortDim =  $\text{argmax}(x, y, z \text{ ranges})$ 
4:   Split data: dataL, dataR
5:   if  $|dataL| > 0$  then
6:     Create tree.L
7:     buildTree_serial( dataL, tree.L, ... )
8:   end if
9:   if  $|dataR| > 0$  then
10:    Create tree.R
11:    buildTree_serial( dataR, tree.R, ... )
12:   end if
13: else
14:   Store data (a single point)
15: end if
```

Building the tree

`buildTree_parallel` performs ORB using a multiple compute nodes

Algorithm 4: `buildTree_parallel(data, tree, comm, ...)`

- 1: Call `getSortDim(...)`: calculates x, y, z mins, maxs, ranges, partition center, and returns *sortDim*
 - 2: Sort *data* over *sortDim* using `parallelSort(data, sortDim, comm, ...)`
 - 3: **if** *myRank* < *numNodes*/2 **then**
 - 4: Create *tree.L*, *commL*
 - 5: `buildTree_parallel(data, tree.L, comm, ...)`
 - 6: **else**
 - 7: Create *tree.R*, *commR*
 - 8: `buildTree_parallel(data, tree.R, comm, ...)`
 - 9: **end if**
-

It is assumed that $tree.n > 1$ will never occur in `build/tree_parallel` since we usually deal with large amounts of data

Building the tree

`getSortDim` finds the longest axis and stores several key tree fields

Algorithm 5: `getSortDim(data, tree, comm, ...)`

- 1: Each process gets its local x, y, z min and max
 - 2: Rank 0 receives these, determines the global x, y, z min and max, determines the `sortDim`, and Bcast's all of these values back to the other nodes
 - 3: The global mins/maxs, partition center, and partition radius are stored in *tree*
 - 4: return *sortDim*
-

Searching the tree

`searchTree_serial` returns the number of points within a given radius about a given point

Algorithm 6: `searchTree_serial(tree, rad, point)`

```
1: found = 0
2:  $d = \sqrt{\sum_{i=1}^3 (point[i] - tree.c[i])^2}$ 
3: if  $d \leq rad + tree.rad$  then
4:   if  $tree.L = NULL \ \&\& \ tree.R = NULL$  then
5:     return 1
6:   else
7:     if  $tree.L \neq NULL$  then
8:       found += searchTree_serial(tree.L, rad, point)
9:     end if
10:    if  $tree.R \neq NULL$  then
11:      found += searchTree_serial(tree.R, rad, point)
12:    end if
13:  end if
14: end if
```

Searching the tree

search501 reads the 501-st data file and loops through the points contained within (as well as the three given radii), calling `searchTree_serial` for each

Algorithm 7: `search501(tree, path, ...)`

1:

Parallel sorting

We had to make several significant alterations to our `parallelSort` program in order to integrate it into our KD tree project

Changes:

- Make rank 0 do work
- Conversion to function
- better `adaptBins`

Making rank 0 do work:

- Initially, rank 0 was just a master node which coordinated the other worker nodes
- this technique is very inefficient for parallel ORB since it requires us to switch to serial mode sooner
- The solution involved 1) cleverly altering a large number of if statements in the code and 2) changing how certain types of sends/recvs were handled

Building the tree

Here is how `parallelSort` is structured now that it is a function

Algorithm 8: `parallelSort(data, rows, myRank, sortDim, comm, ...)`

- 1: Locally sort *data* on each compute node using a qsort
 - 2: Determine the global min/max of the *sortDim*
 - 3: Create linearly spaced bins over range
 - 4: Bin the *data*
 - 5: **while** *uniformity* < *threshold* && *iterations* < *M* **do**
 - 6:
 - 7: **end while**
-

rank 0 multiple communicators
old new alternating

Validation

MATLAB Demos:

- 2D
- 3D

Results

Conclusions